The EU AI Act Readiness Checklist

Twelve checks for founders and small teams who build or use AI in Europe. The Act lands in waves through 2027 — three of them are already binding. Written in plain English, mapped to real articles. Updated October 2026.

The timeline you're working against

1 Aug 2024	AI Act enters into force
2 Feb 2025	Prohibited practices banned + AI literacy duty (Art. 4-5)
2 Aug 2025	GPAI model obligations, governance, penalties provisions
2 Aug 2026	High-risk Annex III rules + Art. 50 transparency + sandboxes — now in force
2 Aug 2027	Annex I regulated-product obligations

1. Inventory every AI system you touch

List every tool that uses AI — chatbots, CV screeners, content generators, analytics copilots — whether you built it or bought it. Split the list into 'we provide it' vs 'we only use it'. You can't comply with what you can't name.

Why: Providers and deployers carry different duties. Most SMBs are deployers, but the moment you rebrand, fine-tune, or substantially modify a system, you become a provider with the full obligations.

Applies: Now — Art. 4 literacy applies since Feb 2025

2. Train your team (AI literacy)

Give everyone who operates an AI system a working briefing: what the system does, what it can't do, and when to escalate. A one-page internal guide per tool is enough — keep it in your wiki.

Why: Article 4 requires AI literacy for anyone dealing with AI systems on your behalf. It has applied since 2 Feb 2025 and regulators already ask for it.

Applies: Applied 2 Feb 2025 — already binding

3. Screen for prohibited practices

Check your stack against the Art. 5 ban list: social scoring, manipulative dark patterns, untargeted face scraping, workplace emotion recognition, biometric categorisation by sensitive traits, and predictive policing. If anything matches, stop using it.

Why: Fines here are the maximum tier — up to EUR 35M or 7% of global turnover. Prohibitions have applied since 2 Feb 2025.

Applies: Applied 2 Feb 2025 — already binding

4. Flag high-risk uses in Annex III

The classic SMB traps are Annex III use cases: CV screening or candidate ranking (employment), credit scoring or insurance risk-pricing, biometric ID, education admissions, and access scoring for essential services. If you deploy one of these, the high-risk chapter applies to you.

Why: High-risk deployers must follow the provider's instructions, keep meaningful human oversight, feed relevant input data, and retain logs. Some deployers also owe a fundamental-rights impact assessment. Annex I is a separate list — safety components of regulated products (machinery, medical devices, toys) — with obligations landing a year later.

Applies: Applied 2 Aug 2026 (Annex III) · 2 Aug 2027 (Annex I products)

5. Keep AI-interaction disclosure visible

Chatbot disclosure is the provider's job under Art. 50(1) — check your vendor ships it. If you deploy a chatbot on your site, support desk, or DMs, keep that disclosure visible near the first message — don't bury it in a footer. For emotion-recognition or biometric-categorisation tools you deploy, you must inform the people exposed yourself (Art. 50(3)).

Why: Art. 50 splits transparency duties: providers must disclose AI interaction (50(1)) and mark synthetic content (50(2)); deployers must inform people exposed to emotion recognition or biometric categorisation (50(3)) and disclose deep fakes (50(4)).

Applies: Applied 2 Aug 2026 — now binding

6. Disclose deep fakes and public-interest AI text

If you publish deep fakes (AI-generated or manipulated image, audio, or video of real people or events) or AI-generated text that informs the public on matters of public interest, label it visibly as artificial. Machine-readable marking of synthetic content sits with the provider — pick AI media tools that already do it.

Why: Art. 50(4) puts deep-fake and public-interest-text disclosure on deployers. Art. 50(2) puts machine-readable synthetic-content marking on providers — a compliant vendor covers that half for you.

Applies: Applied 2 Aug 2026 — now binding

7. Put a human in the loop — on paper

For each AI system, name who reviews output before it reaches a customer and who can switch it off. Write the escalation rule down: a reviewer must be able to override or halt the system, not just watch it.

Why: Art. 26(2) requires deployers to assign human oversight — the control regulators check first. 'A human glances at it sometimes' is not oversight.

Applies: With each high-risk deployment

8. Demand documentation from vendors

Ask every AI vendor for their conformity status and Art. 13 instructions for use before renewal. No instructions, no deployment — their missing homework becomes your compliance risk.

Why: Deployers must use systems 'in accordance with the instructions' — you can't follow instructions you were never given. The provider's Art. 11 technical file is for regulators; what you need is the deployer-facing documentation.

Applies: Before each procurement / renewal

9. Log what the AI did

Keep automatic records of inputs, outputs, and overrides for every high-risk deployment — who approved what, when, and why. Six months minimum is a sensible floor for deployer logs.

Why: Deployers of high-risk systems must retain automatically generated logs under their control. Good logs also win disputes.

Applies: With each high-risk deployment

10. Watch your national enforcement regime

Greece designates its market-surveillance authority and national sandbox under the Act. Track which regulator owns AI oversight and whether the national sandbox is open — it's the cheapest way to test a borderline system safely.

Why: Member states had until Aug 2025 to stand up authorities; sandboxes must exist by Aug 2026. Early sandbox access is free compliance help.

Applies: Rolling — check each quarter

11. If you build on a GPAI model, know the chain

Using an API model internally (Llama, Mistral, commercial APIs) keeps you a deployer — the model provider carries the GPAI duties. But if you put an AI system built on it on the market under your own name or brand, you become that system's provider under Art. 3(3). Fine-tuning or releasing a modified model counts too. Check before you ship.

Why: GPAI obligations applied 2 Aug 2025. An own-name market launch or 'substantial modification' is the line between deployer and provider.

Applies: Applied 2 Aug 2025 — already binding

12. Budget 30 minutes per quarter for a review

Re-run this checklist quarterly: new tools in the inventory, new guidance from the AI Office, any change in what your systems do. The Act phases in through 2027 — readiness is a habit, not a project.

Why: Guidance, codes of practice, and delegated acts keep landing. A quarterly pass keeps you ahead of each new application date.

Applies: Quarterly

Where to check official guidance

- **Regulation (EU) 2024/1689** — the Act itself — EUR-Lex, consolidated text
- **AI Act Service Desk** — the Commission's official implementation guidance
- **National sandbox (GR)** — test a borderline system with regulator cover

This checklist is practical orientation, not legal advice. For borderline classifications, talk to counsel or use your national sandbox before shipping.

Want your AI stack audited against the Act? **cloudless.gr** maps your systems, classifies the risk, and writes the runbook — starting with a free 30-minute audit: **cloudless.gr/contact**